LITS: An Optimized Learned Index for Strings

Yifan Yang, Shimin Chen
SKLP, ACS, Institute of Computing Technology, CAS
University of Chinese Academy of Sciences

Overview

Background

- existing learned indexes fail to outperform traditional indexes when indexing **string keys**

Characteristics of String keys

- long, variable-sized, with skewed prefixes

string keys are common in real-world

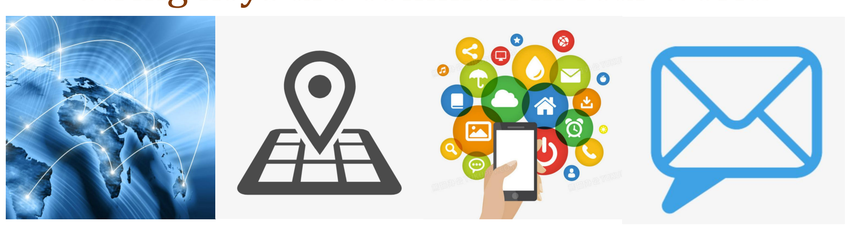

URL address user ID email

Existing Indexes Optimized for Strings

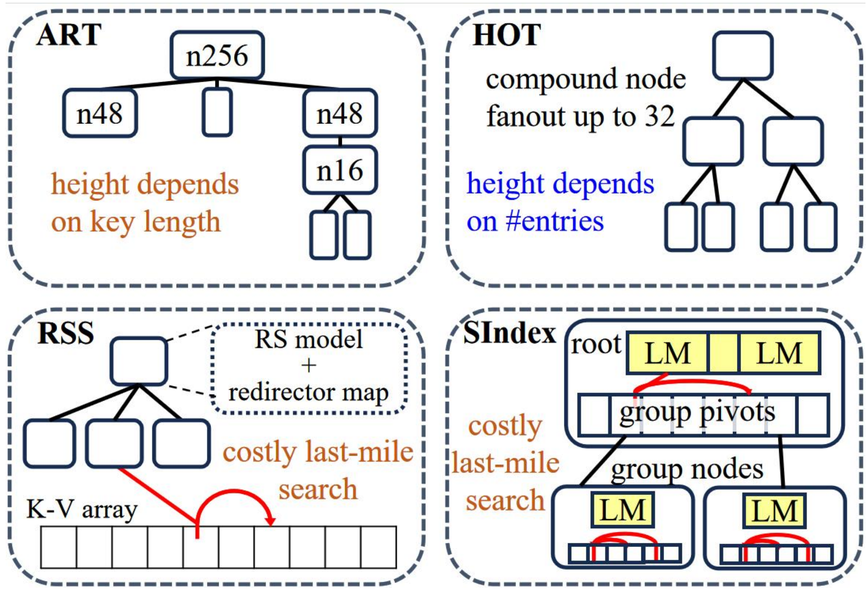

	tree height	node search	last-mile search
ART	×	√	n/a
HOT	--	√	n/a
RSS	√	√	×
SIndex	√	√	×
LITS	√	√	n/a

our solution: LITS

Learned Index with Hash-enhanced Prefix Table and Sub-tries

LITS Structure

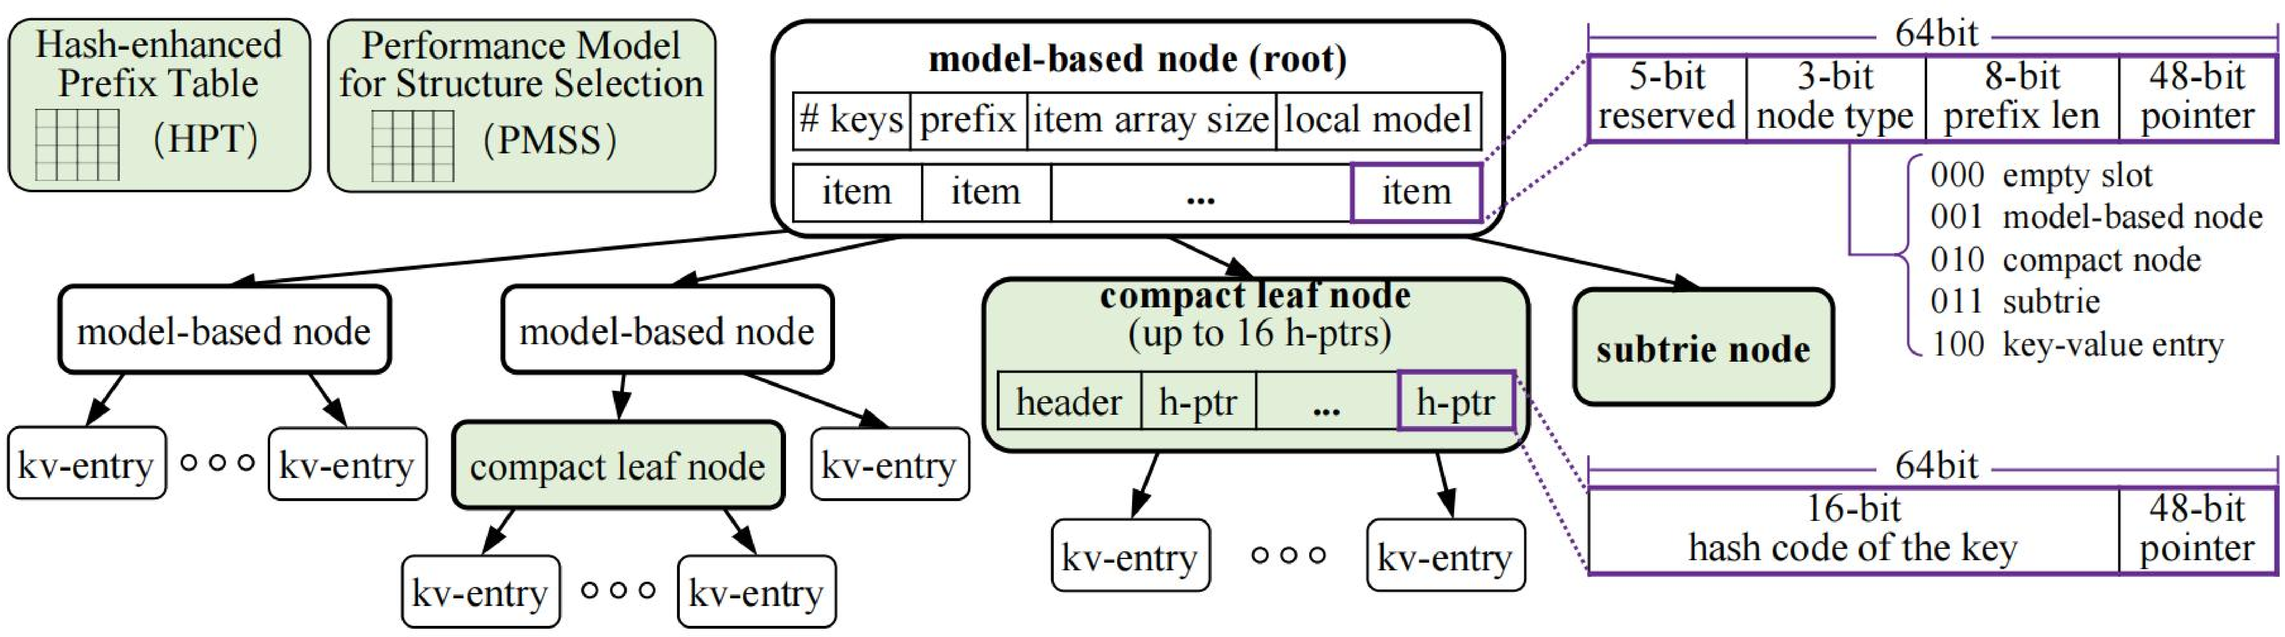

Hash-enhanced Prefix Table

CDF Model in LITS: Hash-enhanced Prefix Table (HPT)

Idea of HPT: learn the pattern of a string data set by better approximating $prob(c|P)$
($prob(c|P)$: conditional probability of the next character being c given the prefix P)

target string: bac

cdf / prob	a	b	c
0	0.1 0.2	0.3 0.4	0.7 0.3
1	0.2 0.1	0.3 0.5	0.8 0.2
2	0.1 0.3	0.4 0.1	0.5 0.5
3	0.4 0.4	0.8 0.1	0.8 0.1

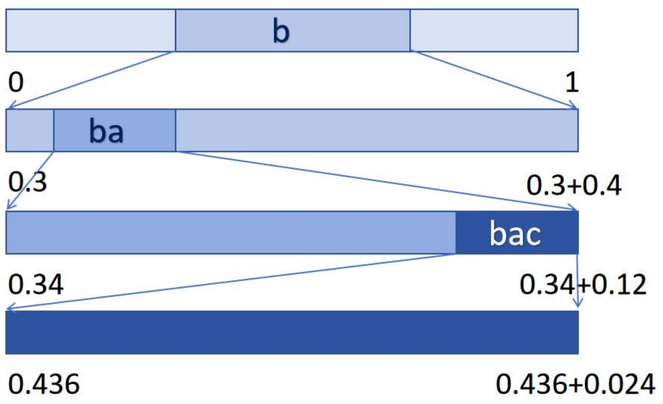

Fig. An illustration of the CDF computation using the HPT for string "bac".

Performance Model for Structure Selection

Hardness of String Data Set: Group Partial Key Length (GPKL)

- GPKL reflects the hardness of modeling a string data set
- GPKL can be computed efficiently by reading the strings in one pass

Idea of PMSS: select the optimal structure for a sub-trie based on the hardness of a subset and offline benchmark tests

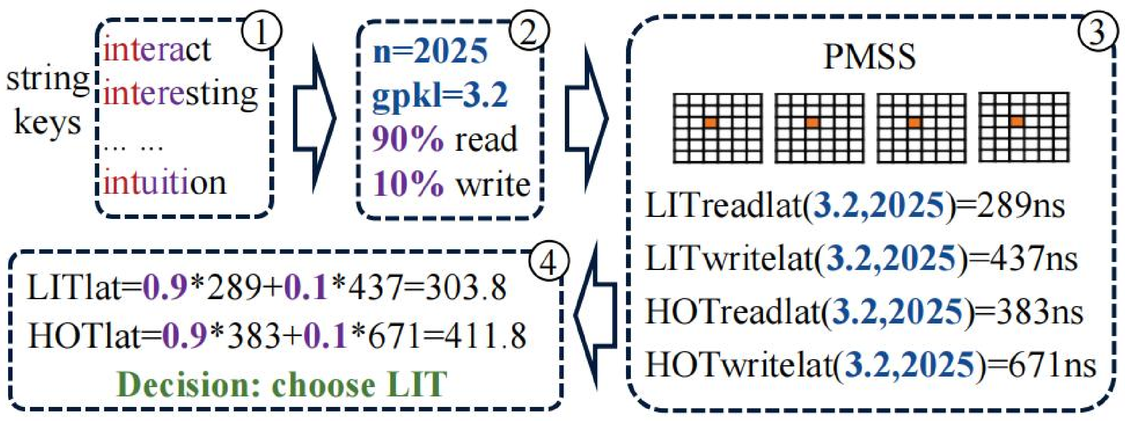

Compact Leaf Node

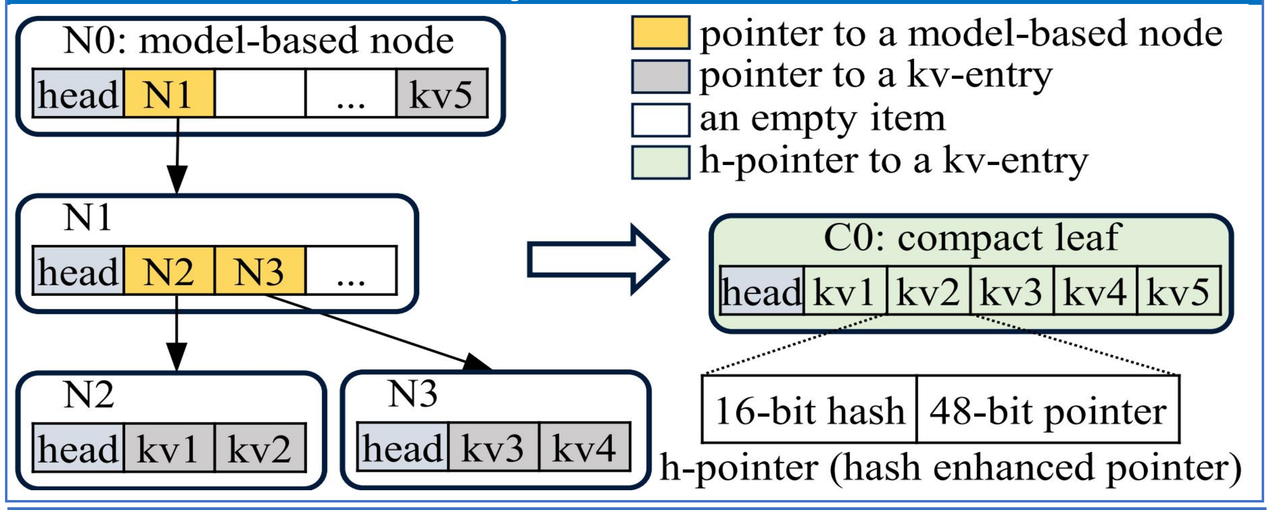

Evaluation

- 7 real-world string data sets and 4 synthetic string data sets
- read-only workload throughput (Mops)

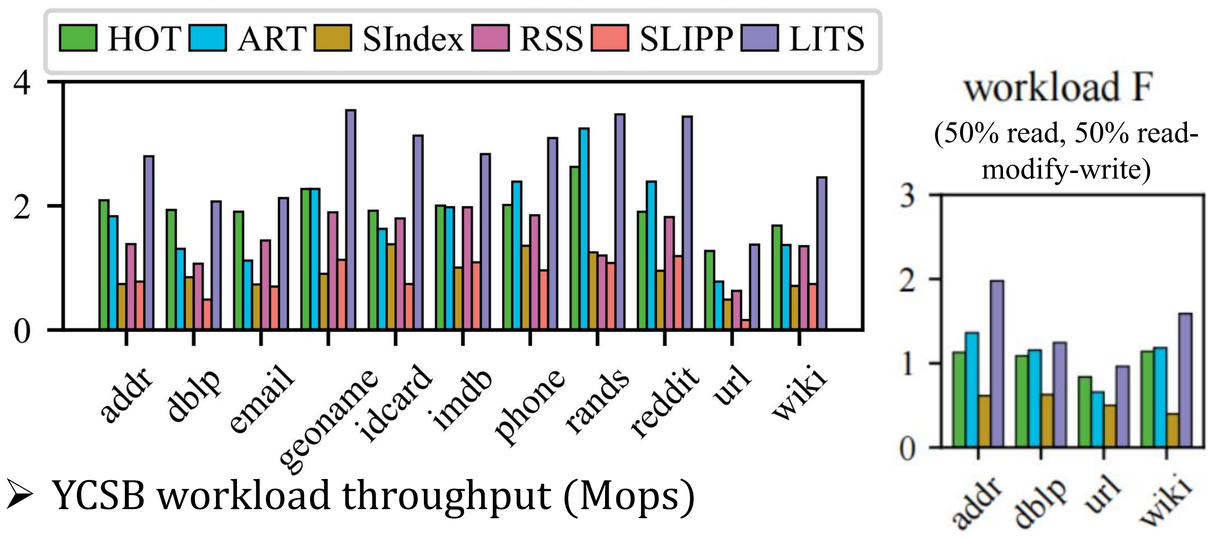

- YCSB workload throughput (Mops)

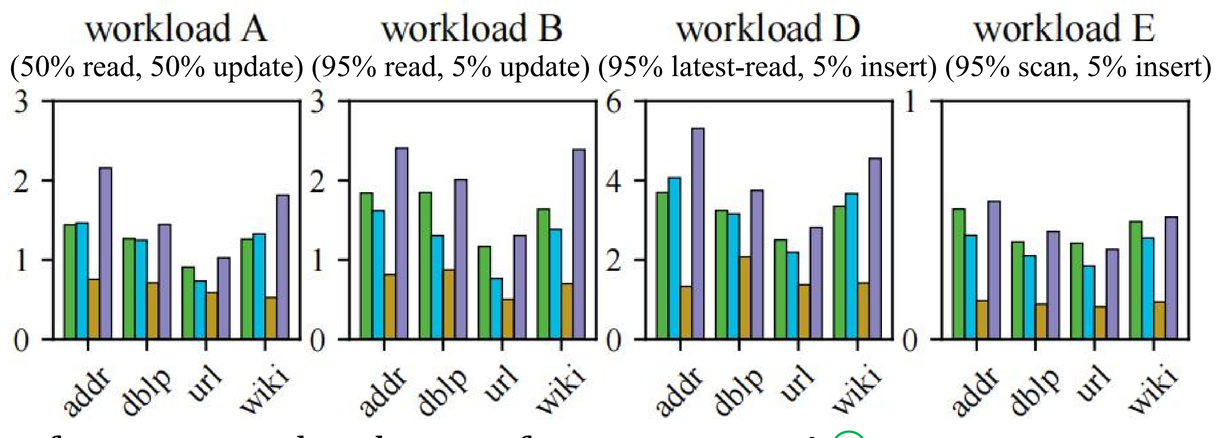

for more results please refer to our paper! ☺

University of Chinese Academy of Sciences
INSTITUTE OF COMPUTING TECHNOLOGY, CHINESE ACADEMY OF SCIENCES